

\$3.75K

Model MAE

\$4.90K

Model RMSE

83 %

Model R2

61 %

Under-Predicted %

\$ 3.53M



predicted_claim

-57.06

Avg Bias

Where are the biggest misses?

policyholder_id	actual_claim	predicted_claim	Abs Error
1166	27,941.28	8,376.44	19,564.84
1151	26,236.58	7,986.48	18,250.10
1146	29,952.82	8,458.41	21,494.41
1187	32,734.19	16,253.11	16,481.09
1110	23,046.37	7,046.39	15,999.98
1336	58,571.07	43,026.70	15,544.37
1119	23,922.02	8,429.41	15,492.61
1101	22,144.03	36,597.81	14,453.79
100789		3,529,323.61	1,001,545.20



Model Performance

The model explains 83% of claim variation ($R^2 = 0.83$) - a strong fit overall. But the errors that remain aren't random; they follow a pattern worth addressing. A small set of policyholders drives most of the error, with individual misses topping \$7K. These are cases the model's current features simply don't capture well. More importantly, the model consistently under-predicts: net signed error sits at - \$15K, and 58% of customers are under-charged. This isn't noise - it's a systematic bias with direct revenue impact. Bottom line: the model is directionally sound but under-prices risk. Priority fix: refine features for high-cost claim segments to close the under-prediction gap and stop the revenue leakage.